

# modeling

April 15, 2025

## 1 Modeling

The goal is to predict the expected value of the total claim amount per exposure unit (year).

- Model the number of claims with a Poisson distribution, and the average claim amount per claim, with a Gamma distribution.

```
[1]: import sys
from pathlib import Path
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import PoissonRegressor, GammaRegressor
from sklearn.model_selection import train_test_split
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import (
    FunctionTransformer,
    OneHotEncoder,
    StandardScaler,
    KBinsDiscretizer,
)

# Add parent directory to sys.path
parent_dir = Path().resolve().parent
sys.path.append(str(parent_dir))

from src.utils import replace_birthdate_with_age, load_data, plot_boxplots
from src.metrics import score_estimator
```

### 1.1 Data load and pre-processing

- load and preprocess data

```
[2]: # load feature data
file_path = parent_dir / 'features.parquet'
df_feat = load_data(file_path)
```

```

# preprocess feature data
df_feat = df_feat[df_feat['Exposure'] > 0.2 ]
#df_feat = df_feat[df_feat['ClaimNb'] < 5]

df_feat = replace_birthdate_with_age(df_feat, 'BirthD',
    ↳reference_date='2023-01-01')
df_feat['VehGas'] = df_feat['VehGas'].fillna('G3')
df_feat["Exposure"] = df_feat["Exposure"].clip(0.1, 1)
df_feat["ClaimNb"] = df_feat["ClaimNb"].clip(upper=4)
df_feat["DriverAge"] = df_feat["DriverAge"].clip(19, 85)
df_feat['VehAge'] = df_feat['VehAge'].clip(0, 20)
df_feat['BonusMalus'] = df_feat['BonusMalus'].clip(0, 100)

# load target data
file_path = parent_dir / 'target.parquet'
df_target = load_data(file_path)

# preprocess target data
df_target = df_target.groupby('IDpol', as_index=False).agg({'ClaimAmount':
    ↳'sum'})
df_target['ClaimAmount'] = df_target['ClaimAmount'].clip(0, 100000)

# merge feature and target data
df_feat["IDpol"] = df_feat["IDpol"].astype(int)
df_feat.set_index("IDpol", inplace=True)
df = pd.merge(df_feat, df_target, on='IDpol', how='left')

#df = df[(df['IDpol'] > 4000000) & (df['IDpol'] < 5000000)]
#df = df[df['ClaimNb'] > 0]

```

DataFrame Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 678013 entries, 0 to 678012

Data columns (total 12 columns):

| # | Column     | Non-Null Count  | Dtype   |
|---|------------|-----------------|---------|
| 0 | IDpol      | 678013 non-null | float64 |
| 1 | ClaimNb    | 678013 non-null | float64 |
| 2 | Exposure   | 678013 non-null | float64 |
| 3 | Area       | 678013 non-null | object  |
| 4 | VehPower   | 678013 non-null | float64 |
| 5 | VehAge     | 678013 non-null | float64 |
| 6 | BonusMalus | 678013 non-null | float64 |
| 7 | VehBrand   | 678013 non-null | object  |
| 8 | VehGas     | 644112 non-null | object  |
| 9 | Density    | 678013 non-null | float64 |

```

10 Region      678013 non-null object
11 BirthD      678013 non-null object
dtypes: float64(7), object(5)
memory usage: 62.1+ MB
None

```

```

#####
First 5 Rows:

```

|   | IDpol | ClaimNb | Exposure | Area | VehPower | VehAge | BonusMalus | VehBrand | \ |
|---|-------|---------|----------|------|----------|--------|------------|----------|---|
| 0 | 1.0   | 1.0     | 0.10     | 'D'  | 5.0      | 0.0    | 50.0       | 'B12'    |   |
| 1 | 3.0   | 1.0     | 0.77     | 'D'  | 5.0      | 0.0    | 50.0       | 'B12'    |   |
| 2 | 5.0   | 1.0     | 0.75     | 'B'  | 6.0      | 2.0    | 50.0       | 'B12'    |   |
| 3 | 10.0  | 1.0     | 0.09     | 'B'  | 7.0      | 0.0    | 50.0       | 'B12'    |   |
| 4 | 11.0  | 1.0     | 0.84     | 'B'  | 7.0      | 0.0    | 50.0       | 'B12'    |   |

|   | VehGas  | Density | Region | BirthD     |
|---|---------|---------|--------|------------|
| 0 | None    | 1217.0  | 'R82'  | 1967-05-08 |
| 1 | Regular | 1217.0  | 'R82'  | 1967-12-28 |
| 2 | Diesel  | 54.0    | 'R22'  | 1970-08-13 |
| 3 | Diesel  | 76.0    | 'R72'  | 1976-12-05 |
| 4 | Diesel  | 76.0    | 'R72'  | 1976-02-29 |

```

#####
Summary Statistics:

```

|       | IDpol        | ClaimNb       | Exposure      | VehPower      | \ |
|-------|--------------|---------------|---------------|---------------|---|
| count | 6.780130e+05 | 678013.000000 | 678013.000000 | 678013.000000 |   |
| mean  | 2.621857e+06 | 0.053247      | 0.528750      | 6.454631      |   |
| std   | 1.641783e+06 | 0.240117      | 0.364442      | 2.050906      |   |
| min   | 1.000000e+00 | 0.000000      | 0.002732      | 4.000000      |   |
| 25%   | 1.157951e+06 | 0.000000      | 0.180000      | 5.000000      |   |
| 50%   | 2.272152e+06 | 0.000000      | 0.490000      | 6.000000      |   |
| 75%   | 4.046274e+06 | 0.000000      | 0.990000      | 7.000000      |   |
| max   | 6.114330e+06 | 16.000000     | 2.010000      | 15.000000     |   |

|       | VehAge        | BonusMalus    | Density       |
|-------|---------------|---------------|---------------|
| count | 678013.000000 | 678013.000000 | 678013.000000 |
| mean  | 7.044265      | 59.761502     | 1792.422405   |
| std   | 5.666232      | 15.636658     | 3958.646564   |
| min   | 0.000000      | 50.000000     | 1.000000      |
| 25%   | 2.000000      | 50.000000     | 92.000000     |
| 50%   | 6.000000      | 50.000000     | 393.000000    |
| 75%   | 11.000000     | 64.000000     | 1658.000000   |
| max   | 100.000000    | 230.000000    | 27000.000000  |

```

#####

```

Unique Values Per Column:

|            |        |
|------------|--------|
| IDpol      | 678013 |
| ClaimNb    | 11     |
| Exposure   | 181    |
| Area       | 6      |
| VehPower   | 12     |
| VehAge     | 78     |
| BonusMalus | 115    |
| VehBrand   | 11     |
| VehGas     | 2      |
| Density    | 1607   |
| Region     | 22     |
| BirthD     | 25775  |

dtype: int64

#####

Total Missing Values in DataFrame:

|            |       |
|------------|-------|
| IDpol      | 0     |
| ClaimNb    | 0     |
| Exposure   | 0     |
| Area       | 0     |
| VehPower   | 0     |
| VehAge     | 0     |
| BonusMalus | 0     |
| VehBrand   | 0     |
| VehGas     | 33901 |
| Density    | 0     |
| Region     | 0     |
| BirthD     | 0     |

dtype: int64

#####

DataFrame Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 26639 entries, 0 to 26638

Data columns (total 2 columns):

| # | Column      | Non-Null Count | Dtype   |
|---|-------------|----------------|---------|
| 0 | IDpol       | 26639 non-null | float64 |
| 1 | ClaimAmount | 26639 non-null | float64 |

dtypes: float64(2)

memory usage: 416.4 KB

None

#####

First 5 Rows:

|   | IDpol     | ClaimAmount |
|---|-----------|-------------|
| 0 | 1552.0    | 995.20      |
| 1 | 1010996.0 | 1128.12     |
| 2 | 4024277.0 | 1851.11     |
| 3 | 4007252.0 | 1204.00     |
| 4 | 4046424.0 | 1204.00     |

#####

Summary Statistics:

|       | IDpol        | ClaimAmount  |
|-------|--------------|--------------|
| count | 2.663900e+04 | 2.663900e+04 |
| mean  | 2.279864e+06 | 2.278536e+03 |
| std   | 1.577202e+06 | 2.929748e+04 |
| min   | 1.390000e+02 | 1.000000e+00 |
| 25%   | 1.087642e+06 | 6.868100e+02 |
| 50%   | 2.137413e+06 | 1.172000e+03 |
| 75%   | 3.180162e+06 | 1.228080e+03 |
| max   | 6.113971e+06 | 4.075401e+06 |

#####

Unique Values Per Column:

|             |       |
|-------------|-------|
| IDpol       | 24950 |
| ClaimAmount | 12369 |

dtype: int64

#####

Total Missing Values in DataFrame:

|             |   |
|-------------|---|
| IDpol       | 0 |
| ClaimAmount | 0 |

dtype: int64

#####

## 1.2 feature and target definitions

- transform features

```
[3]: # log-transform the target variable
log_scale_transformer = make_pipeline(
    FunctionTransformer(func=np.log), StandardScaler()
)

# create a column transformer for preprocessing
```

```

column_trans = ColumnTransformer(
    [
        ("onehot_categorical", OneHotEncoder(), ["VehBrand", "VehPower", "VehGas", "Region", "Area"]),
        ("standardized_numeric", StandardScaler(), ["VehAge", "DriverAge", "BonusMalus"]),
        ("log_scaled_numeric", log_scale_transformer, ["Density"]),
    ],
    remainder="drop",
)
transmored_features = column_trans.fit_transform(df)
#####

df["Claim_freq"] = df["ClaimNb"] / df["Exposure"]
df["Avg_claim_amount"] = df["ClaimAmount"] / np.fmax(df["ClaimNb"], 1)

```

### 1.3 Model Claim frequency model

- The number of claims (ClaimNb) is a positive integer (0 included).
- discrete events occurring in a given time interval (Exposure) independent from each other.
- model the Claim frequency ClaimNb / Exposure and use Exposure as offset.

```

[4]: df_train, df_test, X_train, X_test = train_test_split(df, transmored_features,
    random_state=0)

# Fit a Poisson regression model for claim frequency
glm_freq = PoissonRegressor(alpha=1e-3, max_iter=1000)
glm_freq.fit(X_train, df_train["Claim_freq"],
    sample_weight=df_train["Exposure"])

scores = score_estimator(glm_freq, X_train, X_test, df_train, df_test,
    target="Claim_freq", weights="Exposure")
print("Evaluation of PoissonRegressor on target Claim_freq")
print(scores)

```

Evaluation of PoissonRegressor on target Claim\_freq

| subset             | train  | test   |
|--------------------|--------|--------|
| metric             |        |        |
| R-squared score    | 0.0111 | 0.0092 |
| mean abs. error    | 0.1646 | 0.1647 |
| mean squared error | 0.1540 | 0.1532 |

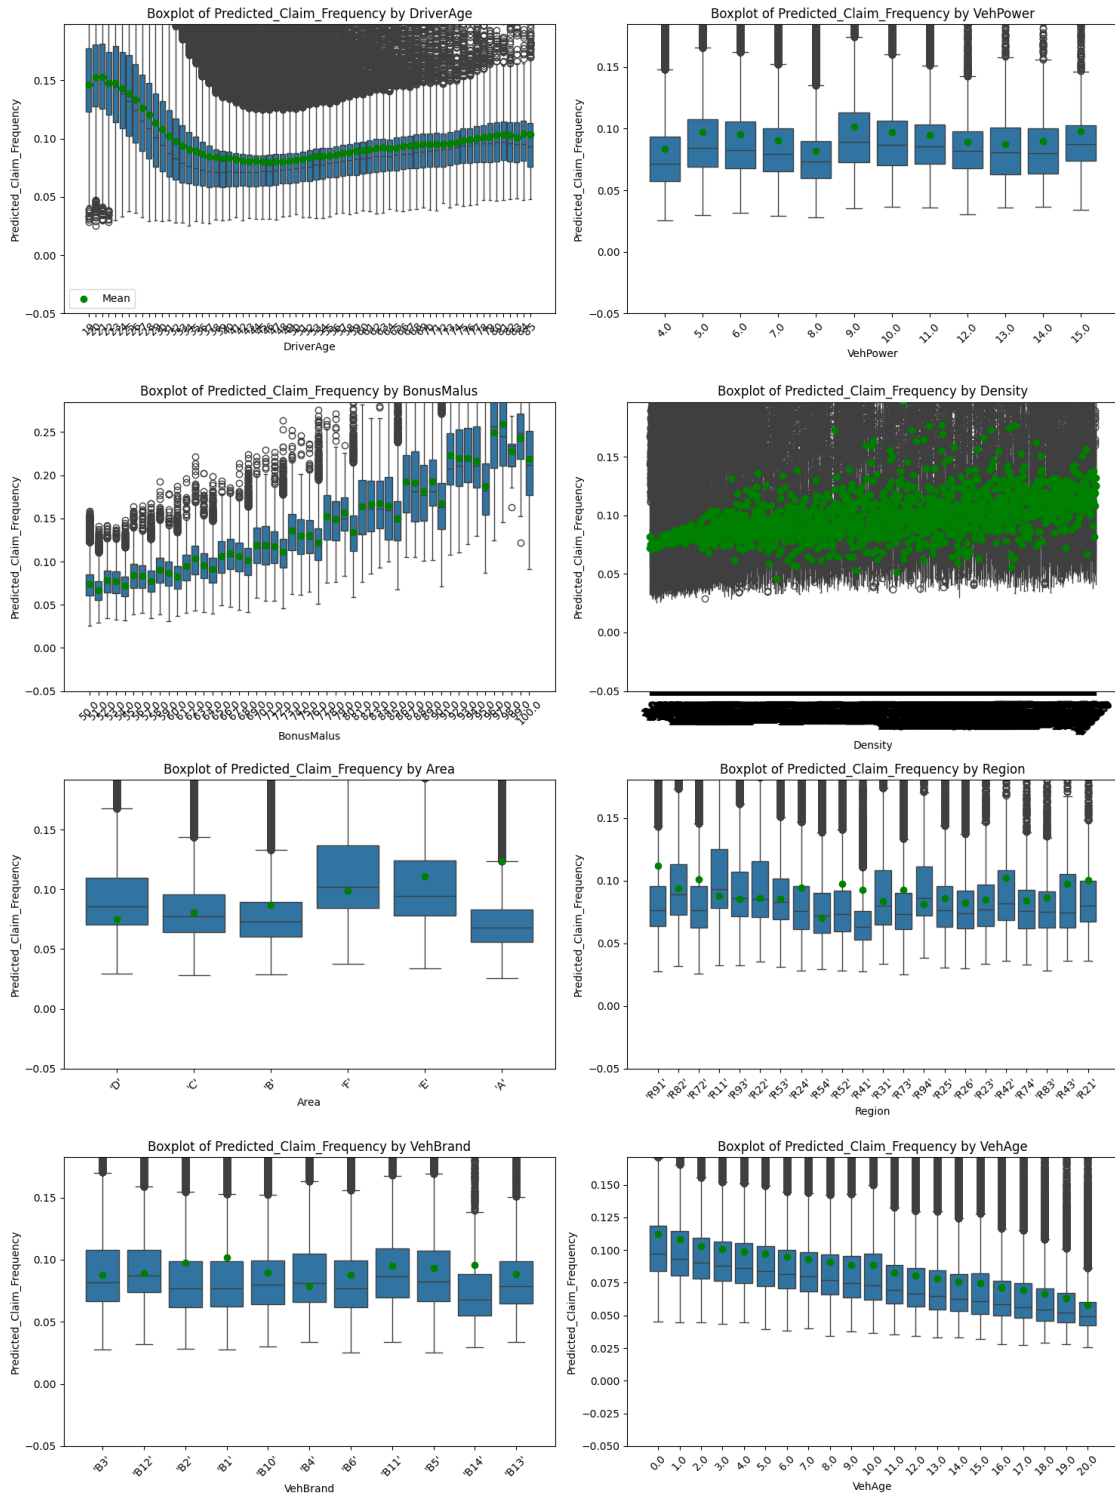
- visualize the predicted values

```

[5]: df_train["Predicted_Claim_Frequency"] = glm_freq.predict(X_train)
    print(df_train.head(10))


```

```
plot_boxplots(['DriverAge', 'VehPower', 'BonusMalus', 'Density', 'Area', 'Region', 'VehBrand', 'VehAge'], 'Predicted_Claim_Frequency', df_train)
```



- visualize the scatter plot of observed and

```
[6]: sns.scatterplot(
    x=df_train['Claim_freq'],
    y=df_train['Predicted_Claim_Frequency'],
    hue=df_train['ClaimNb'],
    s=10 # Set marker size to be smaller
)
plt.xlabel('Observed Claim Frequency')
plt.ylabel('Predicted Claim Frequency')
plt.title('Scatter Plot of Observed vs Predicted Claim Frequency')
plt.xlim(0, 4)
plt.show()
```



#### 1.4 Model average claim amount (Gamma distribution)

- filter out records with 0 claim amount
- use ClaimNb as sample\_weight



```
[7]: mask_train = df_train["ClaimAmount"] > 0
mask_test = df_test["ClaimAmount"] > 0

glm_amount = GammaRegressor(alpha=10.0, max_iter=10000)

glm_amount.fit(X_train[mask_train.values],
               df_train.loc[mask_train, "Avg_claim_amount"],
               sample_weight=df_train.loc[mask_train, "ClaimNb"],
               )

scores = score_estimator(
    glm_amount,
    X_train[mask_train.values],
    X_test[mask_test.values],
    df_train[mask_train],
    df_test[mask_test],
    target="Avg_claim_amount",
    weights="ClaimNb",
)
print("Evaluation of GammaRegressor on target AvgClaimAmount")
print(scores)
```

```
Evaluation of GammaRegressor on target AvgClaimAmount
subset          train          test
metric
R-squared score    1.000000e-04 -1.300000e-03
mean abs. error    1.474881e+03  1.327977e+03
mean squared error  2.590239e+07  1.753600e+07
```

```
[8]: print(
    "actual average claim Amount :          %.2f"
    % df_train["Avg_claim_amount"][df_train["Avg_claim_amount"] > 0].mean()
)
print(
    "Predicted average claim Amount:          %.2f"
    % glm_amount.predict(X_train).mean()
)
```

```
actual average claim Amount :          1774.01
Predicted average claim Amount:          1799.16
```